

Bugzilla-Severity-Detection-Research

June 4, 2018

1 Bugzilla Severity Detection

1.1 Brief introduction about bug report status

RESOLVED

****RESOLVED****

A resolution has been taken, and it is awaiting verification by QA. From here bugs are either re-opened and become REOPENED, are marked VERIFIED, or are closed for good and marked CLOSED.SED.

VERIFIED

QA has looked at the bug and the resolution and agrees that the appropriate resolution has been taken. Bugs remain in this state until the product they were reported against actually ships, at which point they become CLOSED.

CLOSED

The bug is considered dead, the resolution is correct. Any zombie bugs who choose to walk the earth again must do so by becoming REOPENED.

1.2 Imports

```
In [18]: import csv
import re
import json
import numpy as np
import os
import pandas as pd
import sys
import warnings
import re
```

```

import itertools

from sklearn.manifold import *
from sklearn.svm import *
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA
from gensim import *
import nltk
from sklearn.linear_model import LogisticRegression
from nltk.tokenize import *

from sklearn.neural_network import MLPClassifier

from sklearn import preprocessing
from collections import defaultdict
import matplotlib.pyplot as plt
from sklearn import svm
from sklearn.feature_extraction.text import *
from sklearn.metrics import *
from sklearn.naive_bayes import *
from sklearn.model_selection import *
%matplotlib inline
warnings.filterwarnings("ignore")

```

1.3 Retrieving data from preprocessed csv files

```

In [19]: verified = pd.read_csv('summaryList.csv', error_bad_lines=False, quotechar="'", encoding='utf-8')
print("Total number of sentences: ", verified.size)
verified.head(5)

```

Total number of sentences: 58430

```

Out[19]:

```

| | summary | severity | status | \ |
|---|--|-------------|----------|---|
| 0 | [regression] all font-weight are displayed as ... | normal | VERIFIED | |
| 1 | getter/setter bytecodes assume number of atoms... | normal | VERIFIED | |
| 2 | JS_Assert(char * s = 0x1012279c "!\flbase[flind... | critical | VERIFIED | |
| 3 | [FIX]Combobox popups don't have the width of t... | minor | VERIFIED | |
| 4 | Add a strict warning for when an object litera... | enhancement | VERIFIED | |

| | assigned_to | bug_id |
|---|-----------------------|--------|
| 0 | masayuki@d-toybox.com | 365613 |
| 1 | igor@mir2.org | 365692 |
| 2 | general@js.bugs | 365716 |
| 3 | bzbarsky@mit.edu | 365837 |
| 4 | mrbkap@mozilla.com | 365869 |

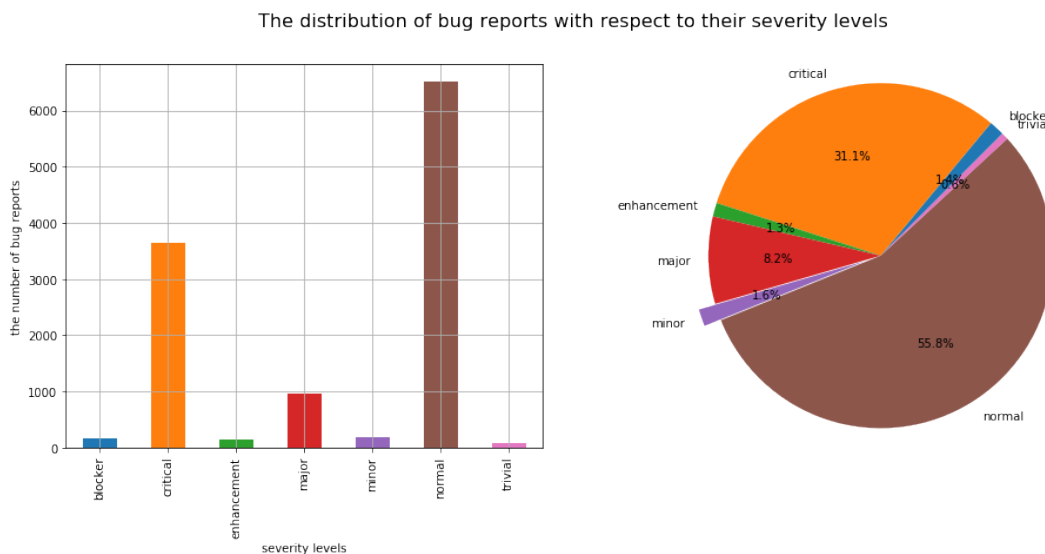
1.4 Multiclass text classification using tf-idf scores of documents

```
In [20]: severity_count = verified.groupby(['severity']).size() # returns a series object
total = severity_count.sum()
severity_percentage = severity_count / total * 100
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 6))
explode = ( 0,0,0,0,0.1, 0, 0)

ax1 = severity_count.plot("bar", ax=axes[0])
ax1.set_ylabel("the number of bug reports")
ax1.set_xlabel("severity levels")
ax1.grid(True)

ax2 = severity_percentage.plot("pie", ax=axes[1], explode=explode, autopct='%1.1f%%', s
ax2.set_ylabel("")
ax2.axis("equal")

plt.suptitle('The distribution of bug reports with respect to their severity levels', f
plt.show()
```



```
In [21]: summaries = [verified.values[id][0] for id in range(len(verified.values))]
severities = [verified.values[id][1] for id in range(len(verified.values))]

In [22]: def clean_str(string):
    """
    Tokenization/string cleaning for all datasets except for SST.
    Original taken from https://github.com/yoonkim/CNN_sentence/blob/master/process_data.py
    """
    string = re.sub(r"[^A-Za-z0-9(),!?\\"`]", " ", string)
```

```

string = re.sub(r"'s", " \'s", string)
string = re.sub(r"'ve", " \'ve", string)
string = re.sub(r"n't", " n\'t", string)
string = re.sub(r"'re", " \'re", string)
string = re.sub(r"'d", " \'d", string)
string = re.sub(r"'ll", " \'ll", string)
string = re.sub(r",", " , ", string)
string = re.sub(r"!", " ! ", string)
string = re.sub(r"\(", " \(", string)
string = re.sub(r"\)", " \)", string)
string = re.sub(r"\?", " \?", string)
string = re.sub(r"\s{2,}", " ", string)
return string.strip().lower()
tokenized = [clean_str(s).split() for s in summaries]

```

```

In [23]: vectorizer = TfidfVectorizer(analyzer='word',tokenizer=word_tokenize, stop_words='english')
         tvec_weights = vectorizer.fit_transform(summaries)
         X = tvec_weights.todense()
         Y = np.asarray(severities)

```

```

In [24]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.1)

```

1.4.1 Multinomial Naive Bayes

```

In [25]: MultiNB = MultinomialNB()
         MultiNB.fit(X_train, Y_train)

         scores = cross_val_score(MultiNB, X_test, Y_test, cv=5, n_jobs=-1)

         print(MultiNB)
         print("Accuracy: %0.2f (+/- %0.2f)" % (scores.mean(), scores.std() * 2))

```

```

MultinomialNB(alpha=1.0, class_prior=None, fit_prior=True)
Accuracy: 0.73 (+/- 0.04)

```

```

In [26]: Y_expect = Y_test
         Y_predict = MultiNB.predict(X_test)
         print("Accuracy: %4.2f " % (accuracy_score(Y_expect, Y_predict)))
         print(confusion_matrix(Y_expect, Y_predict))
         print(recall_score(Y_expect, Y_predict, average=None))
         print(precision_score(Y_expect, Y_predict, average=None))

```

Accuracy: 0.75

```

[[ 0  7  0  0  0 15  0]
 [ 0 265  0  0  0 96  0]
 [ 0  0  0  0  0 20  0]
 [ 0  7  0  0  0 85  0]
 [ 0  1  0  0  0 20  0]

```

```

[ 0 36 0 0 0 612 0]
[ 0 0 0 0 0 5 0]]
[ 0. 0.73 0. 0. 0. 0.94 0. ]
[ 0. 0.84 0. 0. 0. 0.72 0. ]

```

1.4.2 Linear SVM

```

In [27]: lsvm = LinearSVC(dual=False, max_iter=10000)
         lsvm.fit(X_train, Y_train)

         scores = cross_val_score(lsvm, X_test, Y_test, cv=5, n_jobs=-1)

         print(lsvm)
         print("Accuracy: %0.2f (+/- %0.2f)" % (scores.mean(), scores.std() * 2))
         print(scores)

```

```

LinearSVC(C=1.0, class_weight=None, dual=False, fit_intercept=True,
          intercept_scaling=1, loss='squared_hinge', max_iter=10000,
          multi_class='ovr', penalty='l2', random_state=None, tol=0.0001,
          verbose=0)
Accuracy: 0.76 (+/- 0.04)
[ 0.79  0.74  0.76  0.76  0.77]

```

```

In [28]: Y_expect = Y_test
         Y_predict = lsvm.predict(X_test)
         print("Accuracy: %4.2f" % (accuracy_score(Y_expect, Y_predict)))
         print(confusion_matrix(Y_expect, Y_predict))
         print(recall_score(Y_expect, Y_predict, average=None))
         print(precision_score(Y_expect, Y_predict, average=None))

```

```

Accuracy: 0.77
[[ 1  7  0  2  0 12  0]
 [ 0 290 0  1  0 70  0]
 [ 0  0  0  0  0 20  0]
 [ 0 11  0  6  0 75  0]
 [ 0  1  0  0  0 19  1]
 [ 0 36  3  9  0 600  0]
 [ 0  0  0  0  0  4  1]]
[ 0.05  0.8  0.  0.07  0.  0.93  0.2 ]
[ 1. 0.84  0.  0.33  0.  0.75  0.5 ]

```

1.5 Classification among severe class

```

In [29]: df3 = verified.copy()
         sev = ['major', 'critical', 'blocker']
         df3 = df3.loc[df3['severity'].isin(sev)]

```

```
df3.loc[df3.severity == 'blocker', 'severity'] = "critical"
df3.head(5)
```

Out [29]:

| | summary | severity | status | \ |
|----|---|----------|----------|---|
| 2 | JS_Assert(char * s = 0x1012279c "!flbase[flind... | critical | VERIFIED | |
| 9 | large script miscompiles | critical | VERIFIED | |
| 10 | compiling long XML filtering predicate hangs | critical | VERIFIED | |
| 15 | Any use of setter functions causes an assertio... | critical | VERIFIED | |
| 17 | [reflow branch] Crash [@ PresShell::ProcessRef... | critical | VERIFIED | |

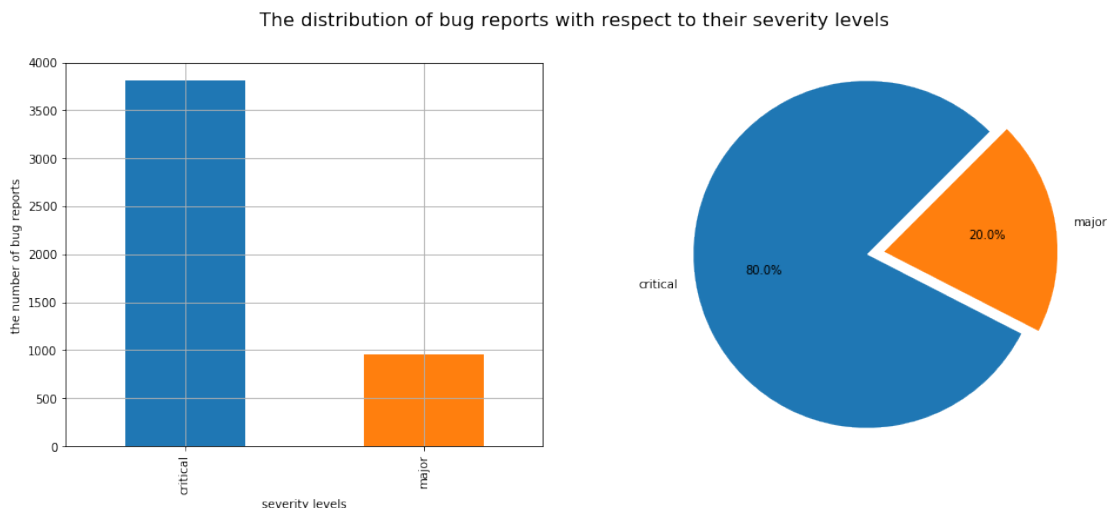
| | assigned_to | bug_id |
|----|---------------------|--------|
| 2 | general@js.bugs | 365716 |
| 9 | igor@mir2.org | 366122 |
| 10 | igor@mir2.org | 366123 |
| 15 | brendan@mozilla.org | 366288 |
| 17 | mats@mozilla.com | 366320 |

```
In [30]: severity_count = df3.groupby(['severity']).size() # returns a series object
total = severity_count.sum()
severity_percentage = severity_count / total * 100
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 6))
explode = (0.1, 0)
```

```
ax1 = severity_count.plot("bar", ax=axes[0])
ax1.set_ylabel("the number of bug reports")
ax1.set_xlabel("severity levels")
ax1.grid(True)
```

```
ax2 = severity_percentage.plot("pie", ax=axes[1], explode=explode, autopct='%1.1f%%', s
ax2.set_ylabel("")
ax2.axis("equal")
```

```
plt.suptitle('The distribution of bug reports with respect to their severity levels', f
plt.show()
```



```

In [31]: summaries = [df3.values[id][0] for id in range(len(df3.values))]
        severities = [df3.values[id][1] for id in range(len(df3.values))]

In [32]: vectorizer = TfidfVectorizer(analyzer='word',tokenizer=word_tokenize, stop_words='engli
        tvec_weights = vectorizer.fit_transform(summaries)
        X = tvec_weights.todense()
        Y = np.asarray(severities)
        X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.1)

In [33]: MultiNB = MultinomialNB()
        MultiNB.fit(X_train, Y_train)

        scores = cross_val_score(MultiNB, X_test, Y_test, cv=5, n_jobs=-1)

        print(MultiNB)
        print("Accuracy: %0.2f (+/- %0.2f)" % (scores.mean(), scores.std() * 2))

        Y_expect = Y_test
        Y_predict = MultiNB.predict(X_test)
        print("Accuracy: %4.2f " % (accuracy_score(Y_expect, Y_predict)))
        print(confusion_matrix(Y_expect, Y_predict))
        print(recall_score(Y_expect, Y_predict, average=None))
        print(precision_score(Y_expect, Y_predict,average=None))

MultinomialNB(alpha=1.0, class_prior=None, fit_prior=True)
Accuracy: 0.82 (+/- 0.01)
Accuracy: 0.85
[[389   3]
 [ 69 15]]
[ 0.99 0.18]
[ 0.85 0.83]

In [34]: lsvm = LinearSVC(dual=False, max_iter=10000)
        lsvm.fit(X_train, Y_train)

        scores = cross_val_score(lsvm, X_test, Y_test, cv=5, n_jobs=-1)

        print(lsvm)
        print("Accuracy: %0.2f (+/- %0.2f)" % (scores.mean(), scores.std() * 2))
        print(scores)

        Y_expect = Y_test
        Y_predict = lsvm.predict(X_test)
        print("Accuracy: %4.2f" % (accuracy_score(Y_expect, Y_predict)))
        print(confusion_matrix(Y_expect, Y_predict))

```

```

print(recall_score(Y_expect, Y_predict, average=None))
print(precision_score(Y_expect, Y_predict, average=None))

LinearSVC(C=1.0, class_weight=None, dual=False, fit_intercept=True,
          intercept_scaling=1, loss='squared_hinge', max_iter=10000,
          multi_class='ovr', penalty='l2', random_state=None, tol=0.0001,
          verbose=0)
Accuracy: 0.84 (+/- 0.04)
[ 0.83  0.82  0.85  0.83  0.87]
Accuracy: 0.89
[[369  23]
 [ 31  53]]
[ 0.94  0.63]
[ 0.92  0.7 ]

```

1.6 Binary class text classification using tf-idf scores of documents

In [35]: `df1 = verified.copy()`

```

df1.loc[df1.severity == 'blocker', 'severity'] = "severe"
df1.loc[df1.severity == 'critical', 'severity'] = "severe"
df1.loc[df1.severity == 'major', 'severity'] = "severe"

df1.loc[df1.severity == 'normal', 'severity'] = "non-severe"
df1.loc[df1.severity == 'minor', 'severity'] = "non-severe"
df1.loc[df1.severity == 'enhancement', 'severity'] = "non-severe"
df1.loc[df1.severity == 'trivial', 'severity'] = "non-severe"

df1.head()

```

Out [35]:

| | summary | severity | status | \ |
|---|---|------------|----------|---|
| 0 | [regression] all font-weight are displayed as ... | non-severe | VERIFIED | |
| 1 | getter/setter bytecodes assume number of atoms... | non-severe | VERIFIED | |
| 2 | JS_Assert(char * s = 0x1012279c "!flbase[flind... | severe | VERIFIED | |
| 3 | [FIX]Combobox popups don't have the width of t... | non-severe | VERIFIED | |
| 4 | Add a strict warning for when an object litera... | non-severe | VERIFIED | |

| | assigned_to | bug_id |
|---|-----------------------|--------|
| 0 | masayuki@d-toybox.com | 365613 |
| 1 | igor@mir2.org | 365692 |
| 2 | general@js.bugs | 365716 |
| 3 | bzbarsky@mit.edu | 365837 |
| 4 | mrbkap@mozilla.com | 365869 |

In [36]: `severity_count = df1.groupby(['severity']).size()` *# returns a series object*
`total = severity_count.sum()`
`severity_percentage = severity_count / total * 100`
`fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 6))`


```

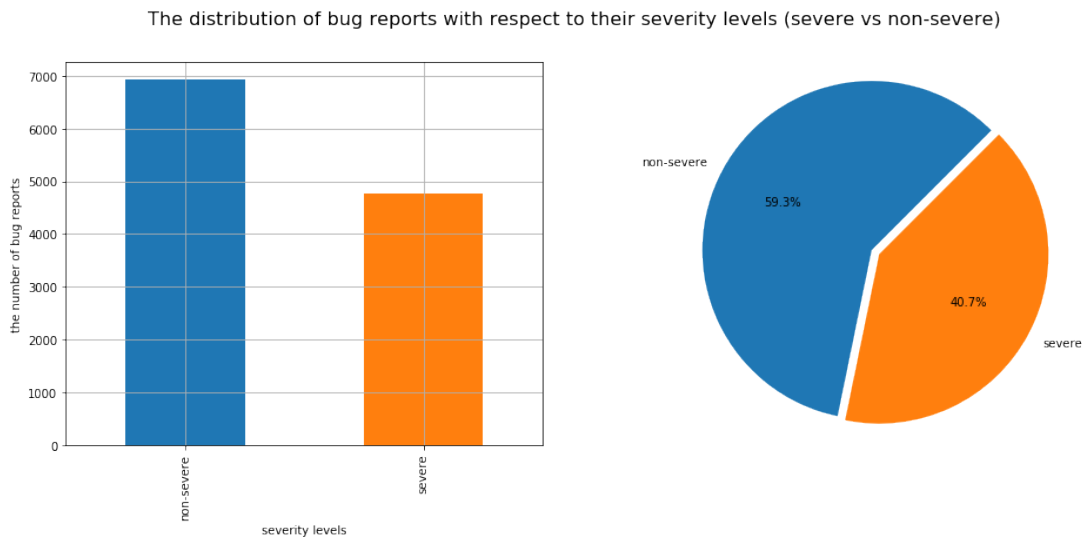
explode = (0, 0.05,)

ax1 = severity_count.plot("bar", ax=axes[0])
ax1.set_ylabel("the number of bug reports")
ax1.set_xlabel("severity levels")
ax1.grid(True)

ax2 = severity_percentage.plot("pie", ax=axes[1], explode=explode, autopct='%1.1f%%', s
ax2.set_ylabel("")
ax2.axis("equal")

plt.suptitle('The distribution of bug reports with respect to their severity levels (se
              fontsize=16)
plt.show()

```



```

In [37]: summaries = [df1.values[id][0] for id in range(len(df1.values))]
          severities = [df1.values[id][1] for id in range(len(df1.values))]

In [38]: vectorizer = TfidfVectorizer(analyzer='word',tokenizer=word_tokenize, stop_words='engli
          tvec_weights = vectorizer.fit_transform(summaries)
          X = tvec_weights.todense()
          Y = np.asarray(severities)

In [39]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.1)

```

1.6.1 Multinomial Naive Bayes

```

In [40]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.1)
          MultiNB = MultinomialNB()
          MultiNB.fit(X_train, Y_train)

```

```

scores = cross_val_score(MultiNB, X_test, Y_test, cv=5, n_jobs=-1)

print(MultiNB)
print("Accuracy: %0.2f (+/- %0.2f)" % (scores.mean(), scores.std() * 2))

MultinomialNB(alpha=1.0, class_prior=None, fit_prior=True)
Accuracy: 0.81 (+/- 0.05)

In [41]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.1)
MultiNB = MultinomialNB()
MultiNB.fit(X_train, Y_train)

scores = cross_val_score(MultiNB, X_test, Y_test, cv=5, n_jobs=-1)

print(MultiNB)
print("Accuracy: %0.2f (+/- %0.2f)" % (scores.mean(), scores.std() * 2))
Y_expect = Y_test
Y_predict = MultiNB.predict(X_test)
print("Accuracy: %4.2f " % (accuracy_score(Y_expect, Y_predict)))
print(confusion_matrix(Y_expect, Y_predict))
print(recall_score(Y_expect, Y_predict, average=None))
print(precision_score(Y_expect, Y_predict, average=None))

MultinomialNB(alpha=1.0, class_prior=None, fit_prior=True)
Accuracy: 0.82 (+/- 0.06)
Accuracy: 0.83
[[652  30]
 [164 323]]
[ 0.96  0.66]
[ 0.8   0.92]

```

```
In [49]: class_names = ["non-severe", "severe"]
```

1.6.2 Linear SVM

```

In [50]: lsvm = LinearSVC(dual=False)
lsvm.fit(X_train, Y_train)

scores = cross_val_score(lsvm, X_test, Y_test, cv=10, n_jobs=-1)

print(lsvm)
print("Accuracy: %0.2f (+/- %0.2f)" % (scores.mean(), scores.std() * 2))

LinearSVC(C=1.0, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, loss='squared_hinge', max_iter=1000,
multi_class='ovr', penalty='l2', random_state=None, tol=0.0001,

```

```
verbose=0)
Accuracy: 0.84 (+/- 0.08)
```

```
In [51]: Y_expect = Y_test
        Y_predict = lsvm.predict(X_test)
        print("Accuracy: %4.2f" % (accuracy_score(Y_expect, Y_predict)))
        print(confusion_matrix(Y_expect, Y_predict))
        print(recall_score(Y_expect, Y_predict, average=None))
        print(precision_score(Y_expect, Y_predict, average=None))
```

```
Accuracy: 0.83
[[627  55]
 [141 346]]
[ 0.92  0.71]
[ 0.82  0.86]
```

```
In [52]: skf = StratifiedKFold(n_splits=5)
        count=0
        total_acc = 0
        M = X
        y = Y

        for train_index, test_index in skf.split(M,y):
            count +=1
            x_train, x_test = M[train_index], M[test_index]
            y_train, y_test = y[train_index], y[test_index]
            lsvm.fit(x_train, y_train)
            y_predict = lsvm.predict(x_test)
            total_acc += accuracy_score(y_test, y_predict)


            # Compute confusion matrix
            cnf_matrix = confusion_matrix(y_test, y_predict)
            np.set_printoptions(precision=2)

            # Plot non-normalized confusion matrix
            plt.figure()
            plot_confusion_matrix(cnf_matrix, classes=class_names,
                                title='Confusion matrix, without normalization')


            plt.show()

        print("Average accuracy: ", total_acc/count)


Confusion matrix, without normalization
[[1256  130]
 [ 342  610]]
```




Confusion matrix, without normalization
[[1183 203]
[311 641]]




Confusion matrix, without normalization
[[1278 107]
[284 668]]



Confusion matrix, without normalization
[[1207 178]
[202 750]]



Confusion matrix, without normalization
[[1239 146]
[309 642]]



Average accuracy: 0.810716897781

In [53]: scores

Out[53]: array([0.81, 0.8 , 0.82, 0.84, 0.85, 0.83, 0.87, 0.91, 0.87, 0.77])

1.6.3 Logistic Regression

```
In [54]: lr = LogisticRegression(dual=False)
         lr.fit(X_train, Y_train)
```

```
scores = cross_val_score(lr, X_test, Y_test, cv=5, n_jobs=-1)
```

```
print(lr)
```

```
print("Accuracy: %0.2f (+/- %0.2f)" % (scores.mean(), scores.std() * 2))
```

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                  intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
                  penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
                  verbose=0, warm_start=False)
```

Accuracy: 0.82 (+/- 0.06)


```
In [55]: Y_expect = Y_test
        Y_predict = lr.predict(X_test)
        print("Accuracy: %4.2f" % (accuracy_score(Y_expect, Y_predict)))
        print(confusion_matrix(Y_expect, Y_predict))
        print(recall_score(Y_expect, Y_predict, average=None))
        print(precision_score(Y_expect, Y_predict, average=None))
```

```
Accuracy: 0.85
[[649  33]
 [148 339]]
[ 0.95  0.7 ]
[ 0.81  0.91]
```

```
In [56]: skf = StratifiedKFold(n_splits=5)
        count=0
        total_acc = 0
        M = X
        y = Y

        for train_index, test_index in skf.split(M,y):
            count +=1
            x_train, x_test = M[train_index], M[test_index]
            y_train, y_test = y[train_index], y[test_index]
            lr.fit(x_train, y_train)
            y_predict = lr.predict(x_test)
            total_acc += accuracy_score(y_test, y_predict)


            # Compute confusion matrix
            cnf_matrix = confusion_matrix(y_test, y_predict)
            np.set_printoptions(precision=2)

            # Plot non-normalized confusion matrix
            plt.figure()
            plot_confusion_matrix(cnf_matrix, classes=class_names,
                                title='Confusion matrix, without normalization')


            plt.show()

        print("Average accuracy: ", total_acc/count)
```


```
Confusion matrix, without normalization
[[1306  80]
 [ 359 593]]
```




Confusion matrix, without normalization
[[1251 135]
[327 625]]




Confusion matrix, without normalization
[[1338 47]
[304 648]]



Confusion matrix, without normalization
[[1273 112]
[221 731]]



Confusion matrix, without normalization
[[1294 91]
[332 619]]



Average accuracy: 0.828173238626

1.7 Feature extraction using gensim

```
In [ ]: def clean_str(string):
        """
        Tokenization/string cleaning for all datasets except for SST.
        Original taken from https://github.com/yoonkim/CNN_sentence/blob/master/process_data.py
        """
        string = re.sub(r"[^A-Za-z0-9(),!?\'`]", " ", string)
        string = re.sub(r"\'s", " \'s", string)
        string = re.sub(r"\'ve", " \'ve", string)
        string = re.sub(r"n\'t", " n\'t", string)
        string = re.sub(r"\'re", " \'re", string)
        string = re.sub(r"\'d", " \'d", string)
        string = re.sub(r"\'ll", " \'ll", string)
        string = re.sub(r",", " , ", string)
        string = re.sub(r"!", " ! ", string)
        string = re.sub(r"\(", " \( ", string)
        string = re.sub(r"\)", " \) ", string)
        string = re.sub(r"?", " ? ", string)
```

```

        string = re.sub(r"\s{2,}", " ", string)
        return string.strip().lower()

In [ ]: class MeanEmbeddingVectorizer(object):
        def __init__(self, word2vec):
            self.word2vec = word2vec
            # if a text is empty we should return a vector of zeros
            # with the same dimensionality as all the other vectors

        def fit(self, X, y):
            return self

        def transform(self, X):
            """
            This method sums all wordvecs of all words in a sentences
            and divides the resulting vector by the len of word count in the sentence
            """
            return np.array([np.sum([self.word2vec[w] for w in words if w in self.word2vec]
                                     [np.zeros(100)], axis=0) / len(words) for words in X ])

In [ ]: sentences = [clean_str(s).split(" ") for s in summaries]
        severities = [df1.values[id][1] for id in range(len(df1.values))]

        model = models.Word2Vec(sentences, size=100, workers=-1, iter=1000)
        words = list(model.wv.vocab)
        model.wv.save_word2vec_format('model.bin')
        # model = models.Word2Vec.load('model.bin')

In [ ]: w2v = dict(zip(model.wv.index2word, model.wv.syn0))
        mev = MeanEmbeddingVectorizer(w2v)
        vec_weights = mev.transform(sentences)

In [ ]: X_train, X_test, Y_train, Y_test = train_test_split(vec_weights, severities, test_size=0

In [ ]: X_train.shape

```

1.7.1 Linear SVM

```

In [ ]: lsvm = LinearSVC()
        lsvm.fit(X_train, Y_train)

        scores = cross_val_score(lsvm, X_test, Y_test, cv=5, n_jobs=-1)

        print(lsvm)
        print("Accuracy: %0.2f (+/- %0.2f)" % (scores.mean(), scores.std() * 2))

In [ ]: Y_expect = Y_test
        Y_predict = lsvm.predict(X_test)
        print("Accuracy: %4.2f" % (accuracy_score(Y_expect, Y_predict)))

```

```

print(confusion_matrix(Y_expect, Y_predict))
print(recall_score(Y_expect, Y_predict, average=None))
print(precision_score(Y_expect, Y_predict, average=None))

```

1.7.2 Logistic Regression

```

In [ ]: lr = LogisticRegression(dual=False)
        lr.fit(X_train, Y_train)

scores = cross_val_score(lr, X_test, Y_test, cv=5, n_jobs=-1)

print(lr)
print("Accuracy: %0.2f (+/- %0.2f)" % (scores.mean(), scores.std() * 2))

```

```

In [ ]: Y_expect = Y_test
        Y_predict = lr.predict(X_test)
        print("Accuracy: %4.2f" % (accuracy_score(Y_expect, Y_predict)))
        print(confusion_matrix(Y_expect, Y_predict))
        print(recall_score(Y_expect, Y_predict, average=None))
        print(precision_score(Y_expect, Y_predict, average=None))

```

```

In [ ]: le = preprocessing.LabelEncoder()
        y = le.fit_transform(Y_train)

pca = PCA(n_components=2)
X_r = pca.fit(X_train).transform(X_train)
plt.figure()
colors = ['red', 'yellow']
lw = 2
target_names = np.array(["non-severe", "severe"])
for color, i, target_name in zip(colors, [0, 1], target_names):
    plt.scatter(X_r[y == i, 0], X_r[y == i, 1], color=color, alpha=.8, lw=lw, label=target_name)
plt.legend(loc='best', shadow=False, scatterpoints=1)
plt.title('PCA of features')

plt.show()

```

1.8 Iteratively checking every classes

```

In [ ]: def check_iteratively(class1, class2):
        df4 = df.copy()
        # df4 = pd.read_csv('summaryListResolved.csv', error_bad_lines=False, quotechar="'")
        df4 = df4.loc[(df4['severity'] == class1) | (df4['severity'] == class2)]
        print(class1.upper() + ' + ' + class2.upper())
        print("*****")
        print(df4.groupby(['severity']).size())

        summaries = [df4.values[id][0] for id in range(len(df4.values))]
        severities = [df4.values[id][1] for id in range(len(df4.values))]

```



```

vectorizer = TfidfVectorizer(analyzer='word',tokenizer=word_tokenize, stop_words='en
tvec_weights = vectorizer.fit_transform(summaries)
X = tvec_weights.todense()
Y = np.asarray(severities)
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.1)

lsvm = LinearSVC(dual=False, max_iter=10000)
lsvm.fit(X_train, Y_train)

scores = cross_val_score(lsvm, X_test, Y_test, cv=5, n_jobs=-1)

print(lsvm)
print("Accuracy: %0.2f (+/- %0.2f)" % (scores.mean(), scores.std() * 2))
print(scores)
Y_expect = Y_test
Y_predict = lsvm.predict(X_test)
print("Accuracy: %4.2f" % (accuracy_score(Y_expect, Y_predict)))
print(confusion_matrix(Y_expect, Y_predict))
print("recall_score: ", recall_score(Y_expect, Y_predict, average=None))
print("precision_score: "precision_score(Y_expect, Y_predict,average=None))

```

```

In [ ]: check_iteratively('enhancement', 'trivial')
        check_iteratively('trivial', 'minor')
        check_iteratively('minor', 'normal')
        check_iteratively('normal', 'major')
        check_iteratively('major', 'critical')
        check_iteratively('critical', 'blocker')

```

```

In [ ]: check_iteratively('enhancement', 'trivial')
        check_iteratively('trivial', 'minor')
        check_iteratively('minor', 'normal')
        check_iteratively('normal', 'major')
        check_iteratively('major', 'critical')
        check_iteratively('critical', 'blocker')

```

1.9 Using word vectors as features

```

In [57]: from sentiment import *

```

```

In [58]: df5 = verified.copy()
        summaries = [df5.values[id][0] for id in range(len(df5.values))]
        severities = [df5.values[id][1] for id in range(len(df5.values))]

        summary_arr = summaries
        summary_arr = convert_tolower(summary_arr)
        summary_arr = remove_punctuation(summary_arr)
        tokenized = tokenize_sentences(summary_arr)

```

```

# documents = remove_stopwords(summary_arr, tokenized)
# stemmed = stem_words(summary_arr, documents)

```

```
In [59]: stemmed = tokenized
```

```
In [60]: class MeanEmbeddingVectorizer(object):
    def __init__(self, word2vec):
        self.word2vec = word2vec
        # if a text is empty we should return a vector of zeros
        # with the same dimensionality as all the other vectors

    def fit(self, X, y):
        return self

    def transform(self, X):
        """
        This method sums all wordvecs of all words in a sentences
        and divides the resulting vector by the len of word count in the sentence
        """
        return np.array([np.sum([self.word2vec[w] for w in words if w in self.word2vec]
                                [np.zeros(100)], axis=0) / len(words) for words in X ]])

```

```
In [62]: sentences = stemmed
```

```

model = models.Word2Vec(sentences, size=100, workers=-1, iter=1000)
words = list(model.wv.vocab)
model.wv.save_word2vec_format('model.bin')
# model = models.Word2Vec.load('model.bin')

```

```

w2v = dict(zip(model.wv.index2word, model.wv.syn0))
mev = MeanEmbeddingVectorizer(w2v)

```

```
X_train, X_test, Y_train, Y_test = train_test_split(vec_weights, severities, test_size=
```

```
In [63]: vec_weights
```

```

Out[63]: array([[ -1.88e-04,  -1.65e-05,  -8.55e-04, ...,  -5.69e-04,  -6.03e-04,
                   1.20e-03],
 [ -1.91e-05,  -7.02e-05,   1.57e-04, ...,   6.53e-04,   8.13e-04,
                   -5.07e-04],
 [ -3.15e-04,   9.49e-04,   7.78e-04, ...,   6.82e-04,   9.34e-04,
                   5.19e-04],
 ...,
 [  1.39e-03,  -9.60e-05,  -1.21e-03, ...,   7.07e-04,   1.49e-03,
                   5.80e-05],
 [  2.03e-04,  -8.10e-04,   1.84e-04, ...,  -2.93e-04,  -2.00e-04,
                   -7.95e-04],
 [ -9.65e-04,   1.41e-03,  -7.03e-04, ...,  -1.53e-03,   1.16e-04,
                   2.13e-04]])

```

```
In [64]: lsvm = LinearSVC(max_iter=5000)
        lsvm.fit(X_train, Y_train)

        scores = cross_val_score(lsvm, X_test, Y_test, cv=5, n_jobs=-1)

        print(lsvm)
        print("Accuracy: %0.2f (+/- %0.2f)" % (scores.mean(), scores.std() * 2))

        Y_expect = Y_test
        Y_predict = lsvm.predict(X_test)
        print("Accuracy: %4.2f" % (accuracy_score(Y_expect, Y_predict)))
        print(confusion_matrix(Y_expect, Y_predict))
        print(recall_score(Y_expect, Y_predict, average=None))
        print(precision_score(Y_expect, Y_predict, average=None))
```

```
LinearSVC(C=1.0, class_weight=None, dual=True, fit_intercept=True,
          intercept_scaling=1, loss='squared_hinge', max_iter=5000,
          multi_class='ovr', penalty='l2', random_state=None, tol=0.0001,
          verbose=0)
```

```
Accuracy: 0.54 (+/- 0.01)
```

```
Accuracy: 0.58
```

```
[[ 0  1  0  0  0 22  0]
 [ 0 48  0  0  0 324 0]
 [ 0  0  0  0  0 17  0]
 [ 0  0  0  0  0 91  0]
 [ 0  0  0  0  0 20  0]
 [ 0  3  0  0  0 633 0]
 [ 0  0  0  0  0 10  0]]
[ 0.  0.13  0.  0.  0.  1.  0. ]
[ 0.  0.92  0.  0.  0.  0.57  0. ]
```

1.10 Sklearn Neural Networks

```
In [65]: df7 = verified.copy()
```

```
df7.loc[df7.severity == 'blocker', 'severity'] = "severe"
df7.loc[df7.severity == 'critical', 'severity'] = "severe"
df7.loc[df7.severity == 'major', 'severity'] = "severe"
```

```
df7.loc[df7.severity == 'normal', 'severity'] = "normal"
```

```
df7.loc[df7.severity == 'minor', 'severity'] = "non-severe"
df7.loc[df7.severity == 'enhancement', 'severity'] = "non-severe"
df7.loc[df7.severity == 'trivial', 'severity'] = "non-severe"
```

```
In [66]: severity_count = df7.groupby(['severity']).size() # returns a series object
        total = severity_count.sum()
```

```

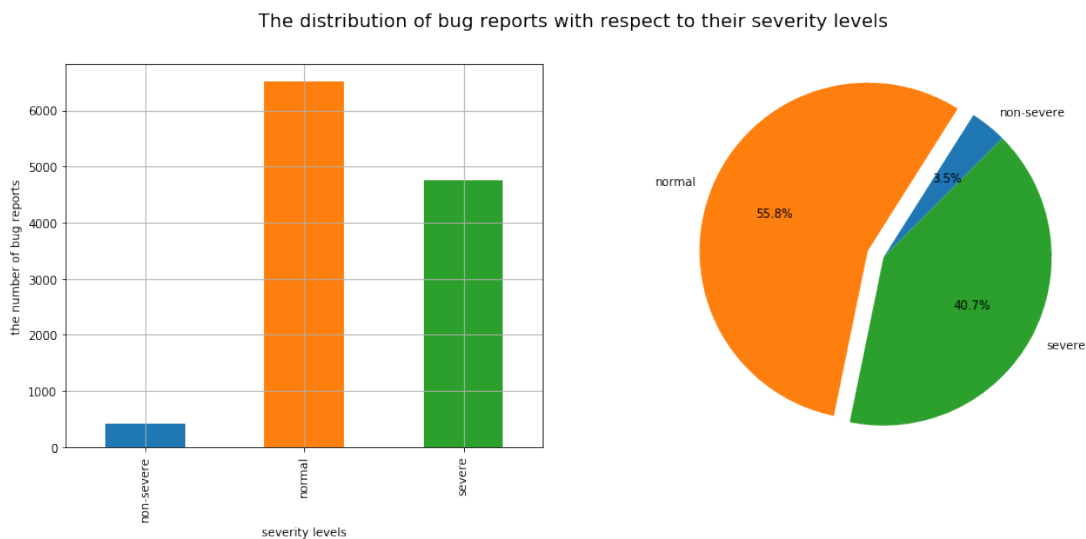
severity_percentage = severity_count / total * 100
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 6))
explode = ( 0,0.1, 0)

ax1 = severity_count.plot("bar", ax=axes[0])
ax1.set_ylabel("the number of bug reports")
ax1.set_xlabel("severity levels")
ax1.grid(True)

ax2 = severity_percentage.plot("pie", ax=axes[1], explode=explode, autopct='%1.1f%%', s
ax2.set_ylabel("")
ax2.axis("equal")

plt.suptitle('The distribution of bug reports with respect to their severity levels', f
plt.show()

```



```

In [67]: summaries = [df7.values[id][0] for id in range(len(df7.values))]
          severities = [df7.values[id][1] for id in range(len(df7.values))]

In [68]: nn = MLPClassifier(hidden_layer_sizes=(100,), activation='tanh', solver='lbfgs', random_
          verbose=True, early_stopping=False, max_iter=1000)

In [69]: class_names = ["non-severe", "normal", "severe"]

In [70]: def plot_confusion_matrix(cm, classes,
          normalize=False,
          title='Confusion matrix',
          cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.

```

```

Normalization can be applied by setting `normalize=True`.
"""
if normalize:
    cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
    print("Normalized confusion matrix")
else:
    print('Confusion matrix, without normalization')

print(cm)

plt.imshow(cm, interpolation='nearest', cmap=cmap)
plt.title(title)
plt.colorbar()
tick_marks = np.arange(len(classes))
plt.xticks(tick_marks, classes, rotation=45)
plt.yticks(tick_marks, classes)

fmt = '.2f' if normalize else 'd'
thresh = cm.max() / 2.
for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
    plt.text(j, i, format(cm[i, j], fmt),
             horizontalalignment="center",
             color="white" if cm[i, j] > thresh else "black")

plt.tight_layout()
plt.ylabel('True label')
plt.xlabel('Predicted label')

```

```

In [71]: skf = StratifiedKFold(n_splits=5)
count=0
total_acc = 0
M = vec_weights
y = severities
y = np.asarray(y)
for train_index, test_index in skf.split(M,y):
    count +=1
    x_train, x_test = M[train_index], M[test_index]
    y_train, y_test = y[train_index], y[test_index]
    nn.fit(x_train, y_train)
    y_predict = nn.predict(x_test)
    total_acc += accuracy_score(y_test, y_predict)

# Compute confusion matrix
cnf_matrix = confusion_matrix(y_test, y_predict)
np.set_printoptions(precision=2)

# Plot non-normalized confusion matrix
plt.figure()

```

```

plot_confusion_matrix(cnf_matrix, classes=class_names,
                      title='Confusion matrix, without normalization')

# Plot normalized confusion matrix
plt.figure()
plot_confusion_matrix(cnf_matrix, classes=class_names, normalize=True,
                      title='Normalized confusion matrix')

plt.show()

print("Average accuracy: ", total_acc/count)

```

Confusion matrix, without normalization

```

[[ 4  71   8]
 [ 3 1151 149]
 [ 0  435 517]]

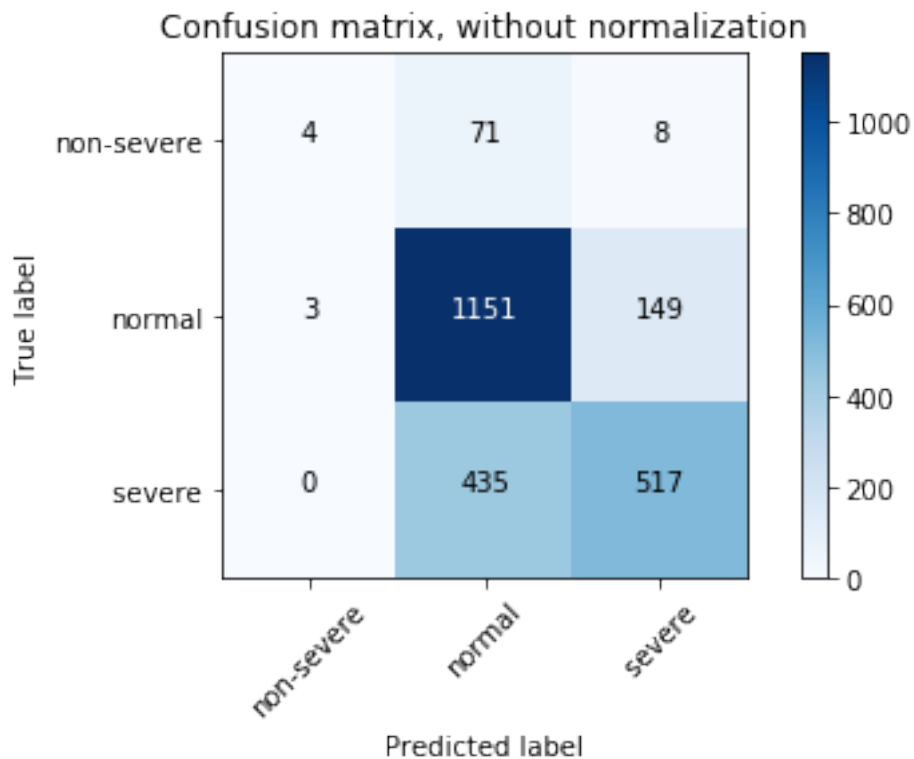
```


Normalized confusion matrix

```

[[ 0.05  0.86  0.1 ]
 [ 0.    0.88  0.11]
 [ 0.    0.46  0.54]]

```







Confusion matrix, without normalization

```
[[ 0  71  12]
 [ 5 1102 196]
 [ 0  399 553]]
```

Normalized confusion matrix

```
[[ 0.    0.86  0.14]
 [ 0.    0.85  0.15]
 [ 0.    0.42  0.58]]
```





Confusion matrix, without normalization

```
[[ 5  68  9]
 [ 4 1182 117]
 [ 2 322 628]]
```

Normalized confusion matrix

```
[[ 0.06  0.83  0.11]
 [ 0.    0.91  0.09]
 [ 0.    0.34  0.66]]
```







Confusion matrix, without normalization

```
[[ 3  76   3]
 [ 3 1138 162]
 [ 0  250 702]]
```

Normalized confusion matrix

```
[[ 0.04  0.93  0.04]
 [ 0.    0.87  0.12]
 [ 0.    0.26  0.74]]
```





Confusion matrix, without normalization

```
[[ 0  78  4]
 [ 0 1186 117]
 [ 0  380 571]]
```

Normalized confusion matrix

```
[[ 0.    0.95  0.05]
 [ 0.    0.91  0.09]
 [ 0.    0.4   0.6 ]]
```





Average accuracy: 0.748081226298